


Your Library, Understood.

A User's Guide to the mcp-calibre Assistant Librarian.



Hand-crafted by jumpifequal.

Give your AI assistant the keys to your library, safely.



Private & Read-Only

Everything runs locally. Calibre databases open strictly read-only (mode=ro). The AI cannot alter or overwrite your files.



Native Comprehension

Equips Claude, ChatGPT, and Codex with 30 distinct tools to search by meaning, read by chapter, extract ISBNs, and view covers.



Millisecond Speed

Bypasses the Calibre GUI. Direct SQLite database connections return search results across a 1,500-book library in under 100 milliseconds.


The Sidecar Architecture keeps your original files perfectly isolated.

The AI Client



The Airgap & Sidecar Cache

MCP Server
Gateway



The only place the server writes data.
Safe to delete anytime.

The Locked Vault



Strictly Read-Only. No calibredb
background process required.

Three distinct search paradigms to pinpoint any concept

1



Metadata (Lexical)

Query Style: Precise filters using Calibre syntax.

Example:

```
tag:security and not  
#letto:true
```

Speed: 3.2-7.8 ms

2



Full-Text (Exact Words)

Query Style: Snippet extraction using BM25 ranking. Finds exact clauses and NEAR groups.

Example:

```
lateral movement
```

Speed: 25-63 ms
(across 10 books)

3



Semantic (Meaning)

Query Style: Multilingual concept matching via local int8 vectors. Asking in Italian finds English passages.

Example:

```
Quali libri spiegano la  
prompt injection?
```

Speed: 75 ms

How the librarian reads any format without losing context.

The Extraction Waterfall

1

Chapter Mapping

Maps the book via TOC or by detecting headings like Chapter 3.

2

Native Extraction

EPUB

Built-in HTML to
Markdown parser

Native Extraction

PDF

PyMuPDF extracts exact
page ranges

Native Extraction

LIT / MOBI


Calibre ebook-convert processes
legacy formats on demand

Character Windows




The AI receives precise, 6,000-character windows centered directly on search hits, preserving vital model tokens.

Making the unreadable searchable in the background.



Visualizing diagrams and covers without spending model tokens.



Auditing and curating the library with read-only diagnostics

Librarian's Digital Desk

Quality Report



Flags metadata anomalies like `raw_filename_title` ('795731065.pdf'), series gaps, or unsorted author names.

Duplicate Detection



Book A



Book B

✓	Field	✓
✓	Title	✓
✓	Author	✓
✓	Year	✓
✓	Edition	✓
✗	Subtitle	⊘ ⓘ


Strict field-by-field matrix checks for true duplication, ignoring subtitles but flagging likely translated versions.

ISBN Discovery




Scans the copyright page of extracted text to instantly discover and validate a missing ISBN-10 or ISBN-13.


Distilling raw chapters into reusable, concept-keyed knowledge.



The Legal Gate: Mechanical evidence of transformation



Millisecond speeds that scale to massive collections



Background Operations Benchmark (1,500-book synthetic library):


- Initial library sync builds a 148 MB index in just 21 seconds.
- Incremental updates take 0.15s.


The Ecosystem at Work: Tracing a complex query

Synthesize agent reliability from my unread security PDFs



Seamlessly connects to your preferred AI environment.





**Your data, your rules,
your AI librarian.**

mcp-calibre bridges the gap between massive personal archives and modern AI—without uploading a single page to the cloud.

Fully open-source. Apache-2.0 Licensed.
Learn more at [jumpifequal/mcp-calibre](https://jumpifequal.com/mcp-calibre)



**Librarian's
Digital Desk**
by Gemini Notebook